

Niraj Patil

Forward Deployed Engineer | Applied AI

+1 447-902-7504 | bnirajpatil9@gmail.com | linkedin.com/in/nirajpatil01 | niraj-portfolio-rho.vercel.app | Chicago, IL (open to relocation, USA)

PROFESSIONAL SUMMARY

Forward-deployed Applied AI engineer who turns ambiguous enterprise problems into production agentic LLM workflows. Built multi-agent RAG systems (MCP orchestration, vector search) and shipped an AI fintech platform from 0 to 1 with confidence-scoring and human-in-the-loop evaluation. Runs technical discovery directly with stakeholders.

CORE COMPETENCIES

Agentic Workflows

RAG / Retrieval

Evaluation Frameworks

LLM Orchestration

Python

Vector Databases

Enterprise Integration

Forward Deployed

WORK EXPERIENCE

Product Engineer May 2026 - Present

Kahana.io | Remote, USA

- Shipped core features for the Oasis AI browser MVP, writing and reviewing full-stack code alongside the engineering team to meet a strict launch deadline.
- Owned the technical integration between product design and engineering, diagnosing and resolving deployment blockers in real time.

Founder & Full-Stack Engineer Jul 2025 - May 2026

Zeno (getzeno.io) | Chicago, IL

- Architected and deployed a full-stack AI fintech platform (Python/Flask, PostgreSQL, AWS S3, React) from 0 to 1, scaling to process 10,000+ monthly transactions in production.
- Built a custom document-parsing pipeline using the Claude API and PostgreSQL schema-mapping to normalize unstructured receipt data with 90%+ accuracy.

Solutions Engineer Aug 2024 - Jul 2025

FACES Consulting, UIUC | Champaign, IL

- Operated inside ambiguous client environments across two AI startups, running structured technical discovery with 20+ stakeholders to surface unstated requirements and translate them into system architecture and integration roadmaps.
- Designed GTM and operational tracking frameworks that drove 15% user adoption within a 90-day pilot window.

Technical Product Manager Oct 2025 - Present

Saayam for All (Nonprofit, Volunteer) | Remote, USA

- Authored use case specs, integration playbooks, and safety guardrails across 15+ service categories, adopted directly by engineering as the implementation guide and as ground-truth data for downstream ML training.

PROJECTS

HawkAI - Autonomous KYC/AML Intelligence Agent

Dual Model Context Protocol (MCP) architecture orchestrating a multi-agent pipeline (SequentialAgent, ResearchAgent, IntelligenceAgent); OFAC SDN screening across 17,557 entries via MongoDB Atlas vector search, generating analyst-grade risk reports in under 90 seconds. Deployed solo on Google Cloud Run. FastAPI, Next.js, Gemini 2.5 Pro, MCP, MongoDB Atlas, Google Cloud Run

AWS-Powered Data Migration & ETL Platform

Production-grade serverless AWS ETL pipeline using Lambda event triggers and Glue (PySpark) to migrate 10,000+ records, automating schema normalization and Parquet conversion to reduce manual analysis time by 80%. AWS S3, Lambda, Glue (PySpark), Athena, Event-Driven Architecture

EDUCATION

M.S., Technology Management | University of Illinois Urbana-Champaign (UIUC) Aug 2025

GPA: 3.93/4.0. Coursework: Data Science, Big Data Infrastructure, Technology Strategy.

B.E., Mechanical Engineering | University of Mumbai Jul 2022

Coursework: Operations Research, Probability & Statistics, Systems Engineering.

SKILLS

Languages & Frameworks: Python (Flask, FastAPI), TypeScript, React, Next.js, Node.js, SQL, REST APIs

AI & Agentic Systems: LLM Integration (Claude, Gemini, Groq), Multi-Agent Orchestration (MCP), RAG, Agentic Systems

Integration & Data: API Integration, Schema Mapping, Data Pipelines, Vector Search (MongoDB Atlas), PostgreSQL, MongoDB

Cloud & DevOps: AWS (S3, Lambda, Glue, Athena), Google Cloud Run, Docker, CI/CD (GitHub Actions)

Customer & Delivery: Technical Discovery, Solution Architecture, Stakeholder Management, Documentation & Playbooks