

Niraj Patil

Applied AI Engineer

+1 447-902-7504 | bnirajpatil9@gmail.com | linkedin.com/in/nirajpatil01 | github.com/Nstar9 | Portfolio | Chicago, IL (open to relocation, USA)

PROFESSIONAL SUMMARY

Applied AI engineer who ships production agents and LLM workflows end-to-end. Built a multi-agent pipeline with MCP orchestration and vector RAG, and an AI fintech platform from idea to production with evaluation and human-in-the-loop monitoring. Strong in Python and TypeScript, and thrives in fast-moving, ambiguous environments where the tools and patterns are still being defined.

SKILLS

AI & Agentic Systems: LLM Integration (Claude, Gemini, Groq), Multi-Agent Orchestration (MCP), RAG, Agentic Systems

Languages & Frameworks: Python (Flask, FastAPI), TypeScript, React, Next.js, Node.js, SQL, REST APIs

Integration & Data: API Integration, Schema Mapping, Data Pipelines, Vector Search (MongoDB Atlas), PostgreSQL, MongoDB

Cloud & DevOps: AWS (S3, Lambda, Glue, Athena), Google Cloud Run, Docker, CI/CD (GitHub Actions)

Customer & Delivery: Technical Discovery, Solution Architecture, Stakeholder Management, Documentation & Playbooks

WORK EXPERIENCE

Product Engineer

May 2026 - Present

Kahana.io | Remote, USA

- Shipped core features for the Oasis AI browser MVP, writing and reviewing full-stack code alongside the engineering team to meet a strict launch deadline.
- Owned the technical integration between product design and engineering, diagnosing and resolving deployment blockers in real time.
- Contributed directly to a #4 Product Hunt launch that acquired 7,000+ users, maintaining system stability with zero critical bugs on launch day.

Technical Product Manager

Oct 2025 - Present

Saayam for All (Nonprofit, Volunteer) | Remote, USA

- Authored detailed use case specs, integration playbooks, and safety guardrails across 15+ service categories, adopted directly by engineering as the implementation guide and as ground-truth data for downstream ML training.
- Led requirements for a volunteer geolocation feature, ran user research with disaster-relief coordinators, and translated findings into Geospatial API integration specs handed directly to engineering.

Founder & Full-Stack Engineer

Jul 2025 - May 2026

Zeno (getzeno.io) | Chicago, IL

- Architected and deployed a full-stack AI fintech platform (Python/Flask, PostgreSQL, AWS S3, React) from 0 to 1, scaling the system to process 10,000+ monthly transactions in production.
- Built a custom document-parsing pipeline using the Claude API and PostgreSQL schema-mapping to normalize unstructured receipt data with 90%+ accuracy.
- Designed a human-in-the-loop validation system and confidence-scoring logic to prevent low-confidence data extractions from reaching 200+ beta users.

Solutions Engineer

Aug 2024 - Jul 2025

FACES Consulting, UIUC | Champaign, IL

- Operated inside ambiguous, undefined client environments across two AI startups, ran structured technical discovery with 20+ stakeholders to surface unstated requirements and translate them directly into system architecture and integration roadmaps.
- Designed GTM and operational tracking frameworks that drove 15% user adoption within a 90-day pilot window.

PROJECTS

HawkAI - Autonomous KYC/AML Intelligence Agent

FastAPI, Next.js, Gemini 2.5 Pro, MCP, MongoDB Atlas, Google Cloud Run

- Implemented a dual Model Context Protocol (MCP) architecture to orchestrate a multi-agent pipeline (SequentialAgent, ResearchAgent, IntelligenceAgent), establishing reliable, bidirectional data flows between external APIs and the core LLM reasoning.
- Integrated OFAC SDN sanctions screening across 17,557 entries via MongoDB Atlas vector search, generating analyst-grade risk reports in under 90 seconds; deployed solo on Google Cloud Run across 39 revisions, entered Google Cloud Agent Hackathon against 10,000+ teams.

AWS-Powered Data Migration & ETL Platform

AWS S3, Lambda, Glue (PySpark), Athena, Event-Driven Architecture

- Built a production-grade serverless AWS ETL pipeline using Lambda event triggers and Glue (PySpark) to migrate 10,000+ records, automating schema normalization and Parquet conversion to reduce manual analysis time by 80%.

Personalized RecSys-LLM - AI Recommendation Engine

Python, Groq, Streamlit, REST API, A/B Testing

- Built a recommendation engine with LLM-driven natural-language onboarding, built-in A/B testing infrastructure, a live public API with Swagger docs, and a Streamlit analytics dashboard.

EDUCATION

M.S., Technology Management | University of Illinois Urbana-Champaign (UIUC)

Aug 2025

GPA: 3.93/4.0. Coursework: Data Science, Big Data Infrastructure, Technology Strategy.

B.E., Mechanical Engineering | University of Mumbai

Jul 2022

Coursework: Operations Research, Probability & Statistics, Systems Engineering.