

Niraj Patil

Applied AI Engineer

+1 447-902-7504 | bnirajpatil9@gmail.com | linkedin.com/in/nirajpatil01 | niraj-portfolio-rho.vercel.app | Chicago, IL (open to relocation, USA)

PROFESSIONAL SUMMARY

Product-minded Applied AI engineer who ships production LLM systems end-to-end. Built agentic pipelines with RAG, tool-calling, and Model Context Protocol (MCP), and took an AI fintech platform from idea to production with real users. Strong on failure modes, observability, and ownership, moving from prototype to production in days.

CORE COMPETENCIES

LLM APIs

RAG / Retrieval

Tool-Calling

AI Agents

MCP

Observability

Python

TypeScript

WORK EXPERIENCE

Product Engineer

May 2026 - Present

Kahana.io | Remote, USA

- Shipped core features for the Oasis AI browser MVP, writing and reviewing full-stack code alongside the engineering team to meet a strict launch deadline.
- Owned the technical integration between product design and engineering, diagnosing and resolving deployment blockers in real time.

Founder & Full-Stack Engineer

Jul 2025 - May 2026

Zeno (getzeno.io) | Chicago, IL

- Architected and deployed a full-stack AI fintech platform (Python/Flask, PostgreSQL, AWS S3, React) from 0 to 1, scaling to process 10,000+ monthly transactions in production.
- Built a custom document-parsing pipeline using the Claude API and PostgreSQL schema-mapping to normalize unstructured receipt data with 90%+ accuracy.

Technical Product Manager

Oct 2025 - Present

Saayam for All (Nonprofit, Volunteer) | Remote, USA

- Authored use case specs, integration playbooks, and safety guardrails across 15+ service categories, adopted directly by engineering as the implementation guide and as ground-truth data for downstream ML training.
- Led requirements for a volunteer geolocation feature, ran user research with disaster-relief coordinators, and translated findings into Geospatial API integration specs handed directly to engineering.

Solutions Engineer

Aug 2024 - Jul 2025

FACES Consulting, UIUC | Champaign, IL

- Operated inside ambiguous client environments across two AI startups, running structured technical discovery with 20+ stakeholders to surface unstated requirements and translate them into system architecture and integration roadmaps.

PROJECTS

HawkAI - Autonomous KYC/AML Intelligence Agent

Dual Model Context Protocol (MCP) architecture orchestrating a multi-agent pipeline (SequentialAgent, ResearchAgent, IntelligenceAgent); OFAC SDN screening across 17,557 entries via MongoDB Atlas vector search, generating analyst-grade risk reports in under 90 seconds. Deployed solo on Google Cloud Run. FastAPI, Next.js, Gemini 2.5 Pro, MCP, MongoDB Atlas, Google Cloud Run

Personalized RecSys-LLM - AI Recommendation Engine

Recommendation engine with LLM-driven natural-language onboarding, built-in A/B testing infrastructure, a live public API with Swagger docs, and a Streamlit analytics dashboard. Python, Groq, Streamlit, REST API, A/B Testing

EDUCATION

M.S., Technology Management | University of Illinois Urbana-Champaign (UIUC)

Aug 2025

GPA: 3.93/4.0. Coursework: Data Science, Big Data Infrastructure, Technology Strategy.

B.E., Mechanical Engineering | University of Mumbai

Jul 2022

Coursework: Operations Research, Probability & Statistics, Systems Engineering.

SKILLS

- Languages & Frameworks:** Python (Flask, FastAPI), TypeScript, React, Next.js, Node.js, SQL, REST APIs
- AI & Agentic Systems:** LLM Integration (Claude, Gemini, Groq), Multi-Agent Orchestration (MCP), RAG, Agentic Systems
- Integration & Data:** API Integration, Schema Mapping, Data Pipelines, Vector Search (MongoDB Atlas), PostgreSQL, MongoDB
- Cloud & DevOps:** AWS (S3, Lambda, Glue, Athena), Google Cloud Run, Docker, CI/CD (GitHub Actions)
- Customer & Delivery:** Technical Discovery, Solution Architecture, Stakeholder Management, Documentation & Playbooks