

{{PHOTO}}

Niraj Patil

+1 447-902-7504 | bnirajpatil9@gmail.com | linkedin.com/in/nirajpatil01 | niraj-portfolio-rho.vercel.app |

Chicago, IL (open to relocation, USA)

PROFESSIONAL SUMMARY

Product-minded Applied AI engineer who ships production LLM systems end-to-end. Built agentic pipelines with RAG, tool-calling, and Model Context Protocol (MCP), and took an AI fintech platform from idea to production with real users. Strong on failure modes, observability, and ownership, comfortable moving from prototype to production in days. Portfolio: niraj-portfolio-rho.vercel.app.

CORE COMPETENCIES

- LLM APIs
- RAG / Retrieval
- Tool-Calling
- AI Agents
- MCP
- Observability
- Python
- TypeScript

WORK EXPERIENCE

Kahana.io May 2026 - Present

Product Engineer / Forward Deployed Engineer | Remote, USA

- Worked directly inside the Oasis AI browser codebase, writing and reviewing feature-level code while translating AI capabilities into technical requirements and resolving deployment blockers as they surfaced.
- Owned the technical handoff between product intent and shipped code for the MVP launch, contributing to a #4 Product Hunt launch with 295 upvotes and 7,000+ trusted users worldwide with zero critical launch-day bugs.
- Acted as the direct technical point of contact between design and engineering, diagnosing integration gaps in real time and driving them to resolution across the full product surface.

Saayam for All Oct 2025 - Present

Product Manager (Technical) | Remote, USA

- Authored detailed use case specs, integration playbooks, and safety guardrails across 15+ service categories, adopted directly by engineering as the implementation guide and as ground-truth data for downstream ML training.
- Led requirements for a volunteer geolocation feature, ran user research with disaster relief coordinators, then translated findings into GeoSpatial API integration specs handed directly to engineering.

Zeno (getzeno.io) Jul 2025 - May 2026

Founder & Full-Stack Engineer | Chicago, IL

- Architected and deployed a full-stack AI fintech platform (Python/Flask, PostgreSQL, AWS S3, React) from 0 to 1, owning the entire lifecycle from technical discovery to processing 10K+ monthly transactions in production.
- Built a document-parsing integration layer (Claude API OCR + LLM) with custom schema-mapping logic in Python/PostgreSQL to standardize unstructured receipt and transaction data into a normalized format, reaching 90%+ extraction accuracy.
- Designed a confidence-scoring and human-in-the-loop validation system to catch low-confidence extractions before they reached production data, maintaining data integrity for 200+ beta users.

FACES Consulting, UIUC Aug 2024 - Jul 2025

Technology Consultant | Champaign, IL

- Operated inside ambiguous client environments across two AI startups, running structured technical discovery with 20+ stakeholders to surface unstated requirements and translate them into system architecture and integration roadmaps.
- Designed GTM and operational tracking frameworks that drove 15% user adoption within a 90-day pilot window.

Strategy Consultant | Champaign, IL

- Built standardized technical evaluation models and reusable integration roadmap templates that accelerated product-market fit validation across 25+ institutional partnerships, supporting a \$5M+ revenue pipeline.
- Delivered implementation roadmaps for three pilot programs, scoping API integration milestones against a 60-day launch deadline.

PROJECTS

HawkAI - Autonomous KYC/AML Intelligence Agent

Implemented a dual Model Context Protocol (MCP) architecture orchestrating a multi-agent pipeline (SequentialAgent, ResearchAgent, IntelligenceAgent) with reliable bidirectional data flows between external APIs and LLM reasoning. Integrated OFAC SDN sanctions screening across 17,557 entries via MongoDB Atlas vector search, generating analyst-grade risk reports in under 90 seconds; deployed solo on Google Cloud Run across 39 revisions.

FastAPI, Next.js, Gemini 2.5 Pro, MCP, MongoDB Atlas, Google Cloud Run

Personalized RecSys-LLM - AI Recommendation Engine

Built a recommendation engine with LLM-driven natural-language onboarding, built-in A/B testing infrastructure, a live public API with Swagger docs, and a Streamlit analytics dashboard.

Python, Groq, Streamlit, REST API, A/B Testing

AlphaForge - ML-Powered Financial Forecasting

Built an ML pipeline predicting next-day equity prices with MAPE of 4.71% (Random Forest) and 5.04% (XGBoost) using chronological train-test splits on OHLCV data; packaged into a Flask demo for non-technical stakeholders.

Python, Random Forest, XGBoost, Flask

EDUCATION

M.S., Technology Management | University of Illinois Urbana-Champaign (UIUC)

Aug 2025

GPA: 3.93/4.0. Coursework: Data Science, Big Data Infrastructure, Technology Strategy.

B.E., Mechanical Engineering | University of Mumbai

Jul 2022

GPA: 3.7/4.0. Coursework: Operations Research, Probability & Statistics, Systems Engineering.

SKILLS

Languages & Frameworks: Python (Flask, FastAPI), TypeScript, React, Next.js, Node.js, SQL, REST APIs

AI & Agentic Systems: LLM Integration (Claude, Gemini, Groq), Multi-Agent Orchestration (MCP), RAG, Agentic Systems

Integration & Data: API Integration, Schema Mapping, Data Pipelines, Vector Search (MongoDB Atlas), PostgreSQL, MongoDB

Cloud & DevOps: AWS (S3, Lambda, Glue, Athena), Google Cloud Run, Docker, CI/CD (GitHub Actions)

Customer & Delivery: Technical Discovery, Solution Architecture, Stakeholder Management, Documentation & Playbooks