

Niraj Patil

Applied AI Engineer

+1 447-902-7504 | bnirajpatil9@gmail.com | linkedin.com/in/nirajpatil01 | github.com/Nstar9 | Portfolio | Chicago, IL (open to relocation, USA)

PROFESSIONAL SUMMARY

Product-minded Applied AI engineer who ships production LLM systems end-to-end. Built agentic pipelines with RAG, tool-calling, and Model Context Protocol (MCP), and took an AI fintech platform from idea to production with real users. Strong on failure modes, observability, and ownership, moving from prototype to production in days. Comfortable across Python and TypeScript, working directly with users.

SKILLS

AI & Agentic Systems: LLM Integration (Claude, Gemini, Groq), Multi-Agent Orchestration (MCP), RAG, Agentic Systems

Languages & Frameworks: Python (Flask, FastAPI), TypeScript, React, Next.js, Node.js, SQL, REST APIs

Integration & Data: API Integration, Schema Mapping, Data Pipelines, Vector Search (MongoDB Atlas), PostgreSQL, MongoDB

Cloud & DevOps: AWS (S3, Lambda, Glue, Athena), Google Cloud Run, Docker, CI/CD (GitHub Actions)

Customer & Delivery: Technical Discovery, Solution Architecture, Stakeholder Management, Documentation & Playbooks

WORK EXPERIENCE

Product Engineer

May 2026 - Present

Kahana.io | Remote, USA

- Shipped core features for the Oasis AI browser MVP, writing and reviewing full-stack code alongside the engineering team to meet a strict launch deadline.
- Owned the technical integration between product design and engineering, diagnosing and resolving deployment blockers in real time.
- Contributed directly to a #4 Product Hunt launch that acquired 7,000+ users, maintaining system stability with zero critical bugs on launch day.

Technical Product Manager

Oct 2025 - Present

Saayam for All (Nonprofit, Volunteer) | Remote, USA

- Authored detailed use case specs, integration playbooks, and safety guardrails across 15+ service categories, adopted directly by engineering as the implementation guide and as ground-truth data for downstream ML training.
- Led requirements for a volunteer geolocation feature, ran user research with disaster-relief coordinators, and translated findings into Geospatial API integration specs handed directly to engineering.

Founder & Full-Stack Engineer

Jul 2025 - May 2026

Zeno (getzeno.io) | Chicago, IL

- Architected and deployed a full-stack AI fintech platform (Python/Flask, PostgreSQL, AWS S3, React) from 0 to 1, scaling the system to process 10,000+ monthly transactions in production.
- Built a custom document-parsing pipeline using the Claude API and PostgreSQL schema-mapping to normalize unstructured receipt data with 90%+ accuracy.
- Designed a human-in-the-loop validation system and confidence-scoring logic to prevent low-confidence data extractions from reaching 200+ beta users.

Solutions Engineer

Aug 2024 - Jul 2025

FACES Consulting, UIUC | Champaign, IL

- Operated inside ambiguous, undefined client environments across two AI startups, ran structured technical discovery with 20+ stakeholders to surface unstated requirements and translate them directly into system architecture and integration roadmaps.
- Designed GTM and operational tracking frameworks that drove 15% user adoption within a 90-day pilot window.

PROJECTS

HawkAI - Autonomous KYC/AML Intelligence Agent

FastAPI, Next.js, Gemini 2.5 Pro, MCP, MongoDB Atlas, Google Cloud Run

- Implemented a dual Model Context Protocol (MCP) architecture to orchestrate a multi-agent pipeline (SequentialAgent, ResearchAgent, IntelligenceAgent), establishing reliable, bidirectional data flows between external APIs and the core LLM reasoning.
- Integrated OFAC SDN sanctions screening across 17,557 entries via MongoDB Atlas vector search, generating analyst-grade risk reports in under 90 seconds; deployed solo on Google Cloud Run across 39 revisions, entered Google Cloud Agent Hackathon against 10,000+ teams.

AWS-Powered Data Migration & ETL Platform

AWS S3, Lambda, Glue (PySpark), Athena, Event-Driven Architecture

- Built a production-grade serverless AWS ETL pipeline using Lambda event triggers and Glue (PySpark) to migrate 10,000+ records, automating schema normalization and Parquet conversion to reduce manual analysis time by 80%.

Personalized RecSys-LLM - AI Recommendation Engine

Python, Groq, Streamlit, REST API, A/B Testing

- Built a recommendation engine with LLM-driven natural-language onboarding, built-in A/B testing infrastructure, a live public API with Swagger docs, and a Streamlit analytics dashboard.

EDUCATION

M.S., Technology Management | University of Illinois Urbana-Champaign (UIUC)

Aug 2025

GPA: 3.93/4.0. Coursework: Data Science, Big Data Infrastructure, Technology Strategy.

B.E., Mechanical Engineering | University of Mumbai

Jul 2022

Coursework: Operations Research, Probability & Statistics, Systems Engineering.